



University of Essex

School of Computer Science and Electrical
Engineering

CAPSTONE PROJECT DISSERTATION

GoLearn — Learning Resource Management System

Bradley Beasley

Supervisor: **Dr Alexandros Voudoris**

April 25, 2024
Colchester

Abstract

There are plenty of services and programs out there for allowing Teachers to upload resources and for their learners to access them. Most of these are not purpose designed for places of learning, such as school, colleges and universities. Those that exist for this specific use-case tend to be over-engineered and can be overly complicated (Moodle is a good example of this). The purpose of GoLearn is to provide a simple LRM experience for teachers and students to interact with the same resources without a large amount of overhead in setting up the system: it should work out of the box.

This report contains a look into methods of implementation for web applications in general, plus the choices in technology and design that were made for GoLearn.

Contents

1	List of Symbols	3
2	Introduction	4
3	Background	5
3.1	Similar Services	5
3.2	Tech Stacks	5
3.3	Databases and ORMs	7
3.4	Testing methods	8
3.5	Authentication	8
4	Technical Specification	10
4.1	Framework	10
4.2	Database and ORM	10
5	Implementation	12
5.1	Data model	12
5.2	Authentication	13
5.3	CRUD Operations	13
5.4	Data Loading	14
5.5	Component design	14
5.6	Emails	14
5.7	Deployment	15
5.8	Known Issues	15
6	Testing	16
6.1	Unit Testing	16
6.2	UAT	16
7	Legal and Ethical	18
8	Project Planning	19
8.1	GIT	20
9	Conclusions	21
9.1	Future Work	21

List of Symbols

API Application Programming Interface

CI/CD Continuous Integration / Continuous Deployment

CRUD Create Read Update Delete

DB Database

GIT Global Information Tracker

JS JavaScript

LAMP Linux, Apache, MySQL, PHP

LEMP Linux, (E)Nginx, MySQL, PHP

LRM Learning Resource Management

MEAN Mongo, Express, Angular, Node

MEVN Mongo, Express, Vue, Node

MERN Mongo, Express, React, Node

NPM Node Package Manager

ORM Object Relational Mapper

OS Operating System

PHP Hypertext Preprocessor

PWA Progressive Web Application

QA Quality Assurance

SEO Search Engine Optimisation

SPA Single Page Application

SQL Structured Query Language

SSR Server Side Rendering

TS TypeScript

UAT User Acceptance Testing

VCS Version Control System

WP WordPress

XHR XMLHttpRequest

Introduction

In the fast-evolving landscape of educational technology, individual passion and innovation can often spark transformative change. The GoLearn project is a testament to this notion, born from the fervent desire of a single individual to streamline the educational experience for both teachers and students. With a deep-seated belief in the power of technology to enhance learning, the creator of GoLearn embarked on a solitary journey to develop a platform that simplifies content organization and sharing.

At its core, GoLearn embodies the creator's unwavering commitment to usability and security. Recognizing the complexities inherent in educational content management, the project sets out to offer a user-friendly interface coupled with robust authentication mechanisms. By leveraging modern technologies and best practices, the goal is to empower educators with a tool that seamlessly integrates into their teaching workflow, freeing them to focus on what matters most: fostering learning.

The journey of building GoLearn is one marked by countless decisions, challenges, and moments of inspiration. From meticulously crafting authentication protocols to architecting data loading strategies, every step reflects the singular vision and dedication of its creator. As a solo venture, the project encapsulates not only technical prowess but also the resilience and determination required to navigate the complexities of software development alone.

As this report delves deeper into the intricacies of GoLearn's development process, it aims to shed light on the triumphs and tribulations experienced along the way. By examining the project's evolution from its humble beginnings to its current state, we gain valuable insights into the transformative potential of individual-driven innovation in the field of educational technology.

Background

3.1 Similar Services

There are various LRMS's available in the market, with some of the more popular one being:

Udemy is a free to access learning site that puts a paywall over all of the courses. Creators can sign up to the site and create a course very quickly. Learners can sign up to courses (paying whatever charge is applied) and have unlimited access to them. They can access these from their own 'Courses' dashboard.

Skillshare is a similar service to Udemy, but is subscription based rather than pay per course. This means that learners can access all of the courses on the site for a monthly fee. Creators can create courses and upload them to the site, with the same access restrictions as Udemy.

LearnDash^[1] is a WordPress plugin that allows you to create and sell courses on your WordPress site. It is a very popular plugin, with over 50,000 active installs and a 4.5-star rating on the WordPress plugin repository. It is a very powerful plugin, with a lot of features, but it is also very complex and can be difficult to use.

Being a plugin for WP, it is very easy to customise how the plugin looks by modifying the theme on the site to best suit the business / individual's needs. The plugin allows teachers to create courses and students to access them, with a lot of customisation options for the courses.

Moodle^[2] is the closest comparison that is currently available to GoLearn. It is a PHP based LMS that is very powerful and feature-rich. It allows learning institutions to design their own courses/modules and add specific 'student' users to them. This is a very similar model to what GoLearn is aiming to achieve. The main difference is that Moodle is very complex and can be difficult to use. It is also very resource-intensive, requiring a lot of server resources to run effectively.

3.2 Tech Stacks

A 'Tech Stack' is the term used to describe the combination of programming languages, frameworks, libraries, and tools that are used to build a software application. The choice of a tech stack is crucial

as it can affect the performance, scalability, and maintainability of the application. The following are some of the popular tech stacks used in web development:

LAMP/LEMP Stacks This stack consists of 4 major components: Linux, Apache/Nginx, MySQL and PHP. Linux is the OS that is used in the stack, with Apache or (E)nginx as the web server, MySQL as the database and PHP as the server-side language. This stack is very popular and is used by many web developers. It is used by various Frameworks (such as Laravel), as well as in vanilla PHP applications (no framework used).

PHP is the programming language of choice for this kind of stack for websites and applications that need to render the pages dynamically, but don't need a lot of reactivity to user input on the page. For example, a blog or a news site would be a good use case for this stack, as they largely need to retrieve data from a db when a page is loaded, and can use caching strategies easily to only make a DB request for cached pages after a certain period after the cached content was stored.

On the other hand, applications that have a lot of user interaction to perform CRUD operations would be harder to write within this stack, as JS would have to be loaded and ran to submit the changed data. Other stacks (as mentioned below) instead rely less on page loads to perform actions, and instead are designed with responsive and reactive implementations in mind.

MEAN/MERN/MEVN Stacks The 'M' refers to the database implementation used (such as mySql, MongoDB)... , the 'E' refers to ExpressJS, which is a NodeJS framework used for all forms of JS applications, and is the layer between the templating framework and NodeJS to serve an application. The 'V', 'A' and 'R' refer to the templating framework used to create the layouts used, and reference 'VueJS', 'AngularJS' and 'ReactJS' respectively. The 'N' refers to 'NodeJS', which is a server-side JS runtime environment that is used to run JS code on the server.

This stack is very popular for creating SPAs and PWAs. These are applications that are designed to be very reactive to user input, and can update the page without needing to reload the page. This is done by using JS to update the page, rather than relying on the server to render the page and send it to the client. This is achieved by requesting information from the server about what to load and how to load it, and then using JS to render the page based on the information received. This is a very powerful stack, but can be difficult to learn and use effectively.

Full Stack Frameworks A full stack framework is a framework that is designed to be used for both the front-end and back-end of an application. This means that the framework can be used to create the server-side code, as well as the client-side code. This is very useful as it means that development teams can rely on a single codebase for the entire application, rather than having to use multiple codebases. This can make development faster and easier. As well as this, it allows type safety and code completion to be used across the entire application easier than in separate front- and back-end applications.

Some examples of full stack frameworks are NextJS and NuxtJS. These frameworks are designed to be used for both the front-end and back-end of an application. They are very powerful and can be used to create very complex applications. They each use a different templating framework (NextJS uses ReactJS, and NuxtJS uses VueJS). The benefit of these frameworks is that they allow

for server-side rendering as well as client-side rendering, which can be very useful for SEO and performance reasons.

One drawback of these frameworks is that there can be a steep learning curve to get started with them. As well as this, the codebase can get very complex and the boundaries between client and server code can get blurred, which can make it difficult to debug and maintain the codebase if it is not well organised.

3.3 Databases and ORMs

A database is a collection of data that is stored in some structured format. Databases are used to store data that can be accessed and manipulated by software applications. There are many different types of databases, each with its own strengths and weaknesses, but all fall into one of two categories: relational or non-relational.

Relational Databases[3] store data in tables, with each table containing rows and columns. The columns represent the attributes of the data, while the rows represent individual records. Relational databases use SQL to query and manipulate the data. Some popular relational databases include MySQL, PostgreSQL and SQLite.

This category of database is called relational because it allows for relationships to be defined between the columns in different tables. For example, a database for a school might have a 'students' table and a 'courses' table. The 'students' table might have a column that references the 'courses' table, which would allow for a relationship to be defined between the two tables. This is called a foreign key, and can be used to enforce that a student can only be enrolled in a course that exists in the 'courses' table.

This is a suitable database design for many kinds of applications, such as those that require complex relationships between the data, or that require consistency on the data. A drawback that this can lead to is that the data can be inflexible, as the schema has to be defined before the data can be stored, which can make it problematic to change the data model once the data has been stored.

Non-Relational Databases[4] store data in a more flexible format, such as key-value pairs, documents, or graphs. Non-relational databases are often used when the data is unstructured or when the data model is likely to change frequently. Some popular non-relational databases include MongoDB, Cassandra, and Redis. This kind of database is often referred to as a NoSQL database.

This category of database stores data in an unstructured format, which allows for it to be more flexible and horizontally scalable. This is useful, as it means that the data can be stored in such a way that the load can be distributed across multiple servers, which can help to improve performance and scalability. A drawback that this can lead to is that the data can be inconsistent, as there may be duplicated data across multiple servers, which can lead to issues with data integrity.

This is a suitable database design for many kinds of applications, such as those that require high performance and scalability, or that require the ability to store large amounts of data. As well

as this, applications that require a lot of flexibility in the data model can benefit from using a NoSQL database, as it allows for the change in the data model over time, without having to make adjustments to the existing data.

An ORM is a layer between the application and the database that allows the application to interact with the database using an object-oriented interface. These are used to simplify the process of interacting with the database, and to make it easier to work with the data in the database. They supply an API for developers to be able to perform CRUD operations on the database without having to construct many (if any) queries or statements to do so.

Some examples of ORMs include Prisma, TypeORM, and Sequelize. These ORMs are designed to be used with different databases, and have different features and capabilities. For example, Prisma is designed to be used with PostgreSQL, MySQL, and SQLite, and has a strong focus on type safety and code generation. TypeORM is designed to be used with TypeScript, and has a strong focus on type safety and code generation. Sequelize is designed to be used with MySQL, PostgreSQL, SQLite, and MSSQL, and has a strong focus on performance and scalability.[5]

3.4 Testing methods

Unit testing[6] is a type of testing that is used to test individual units or components of an application.

It is used to ensure that the individual units of an application are working as expected, and that the various functions behave as expected in isolation from each other. Unit testing is a very important part of the development process, as it can help to catch bugs and issues early in the development process, before they become more difficult and expensive to fix. Unit testing can also help to ensure that the application is working as expected, and that there are no bugs or issues that could cause the application to fail.

User acceptance testing[7] is a type of testing that is used to ensure that the application is working as expected from the user's perspective. It is used to ensure that the application is easy to use, and that the user interface is intuitive and user-friendly. UAT is a very important part of the development process, as it can help to ensure that the application is meeting the needs of the users, and that the users are able to use the application effectively. UAT can also help to identify any issues or bugs that may have been missed during the development process, and can help to ensure that the application is working as expected.

3.5 Authentication

Authentication can be defined as: 'Verifying the identity of a user, process or device, often as a prerequisite to allowing access to resources in an information system.'[8]. Simply put, it refers to any processes used to ensure that only those with the right to access a resource or perform functions can do so. This is a very important part of any application, especially those that can be accessed from a remote location (such as via the internet). This is because protecting the data that is managed by the application is a legal requirement (see Chapter 7).

There are several methods of authentication that can be implemented. A few are listed and described below:

Basic Authentication[\[9\]](#) This is the simplest form of authentication, and is often used in conjunction with a username and password. The user is prompted to enter their username and password, which is then sent to the server. The server then checks the username and password against a list of valid credentials, and if they match, the user is granted access to the resource. This is a very simple form of authentication, but it is not very secure, as the username and password are sent in plain text over the network. This means that anyone who is able to intercept the request can read the username and password and use them to access the resource. For this reason, it is not recommended to use basic authentication for anything other than testing purposes.

Token-Based Authentication[\[10\]](#) This is a more secure form of authentication that is often used in conjunction with a username and password. The user is prompted to enter their username and password, which is then sent to the server. The server then checks the username and password against a list of valid credentials, and if they match, the server generates a token and sends it back to the client. The client then includes the token in all subsequent requests to the server, and the server checks the token against a list of valid tokens to ensure that the user is authenticated. This is a more secure form of authentication, as the token is not sent in plain text over the network. This means that even if someone is able to intercept the request, they will not be able to read the token and use it to access the resource.

OAuth[\[11\]](#) is an open standard for access delegation, commonly used as a way for Internet users to grant websites or applications access to their information on other websites but without giving them the passwords. This is done by authorizing the application to access the data on the user's behalf, and then the application is given a token that can be used to access the data. This is a very secure form of authentication, as the user's password is never sent over the network. This means that even if someone is able to intercept the request, they will not be able to read the password and use it to access the resource.

Technical Specification

4.1 Framework

Of the stacks mentioned in Section 3.2, the NextJS framework was chosen for the GoLearn project. This was chosen for the following reasons:

ReactJS[12] ReactJS is a very powerful and popular front-end framework that is used to create SPAs and PWAs. It is very flexible and can be used to create very complex applications. It is also very easy to learn and use, which makes it a good choice for the GoLearn project.

As well as this, it was a framework that the author had used before, meaning that he could get started with the project quickly and easily and wouldn't have to spend a lot of time learning to use a new framework or architect his own solution from scratch. This proved to save a lot of time in the development process, as the author was able to get started on the project and make progress quickly.

Server Actions[13] As of NextJS 14, the framework included a feature called 'Server Actions'. This allowed functions to be shared between server- and client-side code. Whenever a server action was called, a request is sent to the server as a Fetch/XHR request from the client, the server would handle the request and return the result to the client. This would prove to be a very useful feature, as it allows for CRUD operations to be ran.

4.2 Database and ORM

The ORM chosen for the GoLearn project was Prisma. This was chosen for the following reasons:

Type Safety[14] Prisma is designed to be used with TypeScript, and has a strong focus on type safety and code generation. This means that the author could be confident that the code he was writing was correct and that the database schema was being enforced correctly. This was very important for the GoLearn project, as it was a project that was being developed by a single developer, and there was no QA team to review the code.

Code Generation[15] Prisma uses code generation to generate the database schema and the client-side code that is used to interact with the database. This means that the author could be confident

that the code he was writing was correct and that the database schema was being enforced correctly. This was very important for the GoLearn project, as it was a project that was being developed by a single developer, and there was no QA team to review the code.

Prisma can be used with many different databases, including PostgreSQL, MySQL, and SQLite. One of the factors that was considered was the ease of creating and maintaining the database. Originally the author had considered using a service called 'Supabase', which is a service that provides a PostgreSQL database, and is a service that Prisma has a lot of support for.

Prisma is used by defining a schema in a file called 'schema.prisma', which defines the data models that are used in the application, as well as the relationships between the models. Prisma then uses this schema to generate the database schema and the client-side code that is used to interact with the database. Whenever changes are made to the schema, Prisma can be used to update the database schema and the client-side code, which makes it very easy to maintain the database and the codebase.

Later in the project when the application hosting was being considered, it was decided that the author would host the application on Vercel, which is a service that provides hosting for NextJS applications. Vercel has a feature called 'Vercel for Databases', which allows you to create a PostgreSQL database that is hosted on Vercel. It was a relatively simple process to migrate the database from Supabase to Vercel, and the author was able to do this without too much difficulty. The Vercel database is also a PostgreSQL database, which meant that the author could continue to use Prisma with the database without any issues.

Implementation

5.1 Data model

There was a lot of thought put into the data model for the GoLearn project. The data model was designed to be as simple as possible, while still being able to store all of the information that was needed for the application to function correctly. The following models were implemented, with various relations between them:

User The User model was used to store data about a User, such as username, password, role, etc. The role was used to determine what permissions the user had, such as being able to edit content or not.

Student The Student model was used to store data about Students, such as their home and term addresses, enrolled course and modules. These fields were used to determine which modules and course pages they could access. If a student was not enrolled in a course, they would not be able to access the course page, the same being for the modules.

Teacher The Teacher model was used to store data about Teachers, such as their address and the modules they were teaching. These fields were used to determine which modules pages they could access. If a teacher was not teaching a module, then they can not add content to those modules, nor make edits to them.

Course A course has a name, description and a list of modules that are part of the course. The course is the top-level resource in the GoLearn project, with all other resources being children of the course, or children of children, etc. The students that are enrolled in the course are also associated to a course.

Module A module has a name, description, assignments, units, students and teachers. This is an important model, as it acts as a container for students and teachers to interact with the content. Teachers can create and modify the content in a module and it's children, and students can view the content if they are enrolled in the module.

Assignment Assignments are a type of resource that can be created by teachers within a module, which students can then create submissions for. Assignments have a name, description, due

date, and a list of submissions that have been made for the assignment.

Unit A unit is a resource that teachers can create within a module. This is a method for organising the content within a module into smaller, more manageable chunks. Units have a name, description, and a list of resources that are part of the unit. There can also be sections within a unit, which are used to further organise the content within the unit.

Section A section is a resource that acts as a container for Resources within a Unit. Sections have a name, description, and a list of resources that are part of the section. The reason for this inclusion into the model was to allow teachers to better organise content within a unit such that they can group together resources that are related to each other without having to use naming convention to organise it for them.

Resource A resource is a leaf resource that can be created by teachers within a section and/or. This is the lowest level of resource that can be created, and is used to store the content that is displayed to the students. Resources have a name, description, and a content field that is used to store the content that is displayed to the students. As well as this, they can have any number of files attached to them, which can be used to provide additional resources to the students that is more complex than can be expressed by markdown.

5.2 Authentication

The method of authentication that was chosen for the GoLearn project was token-based authentication. This was chosen as it is a relatively simple form of authentication that is secure and easy to implement. Within the database, a 'UserSession' table was created to store the tokens that are generated for each user. When a user logs in, a token is generated and stored in the 'UserSession' table, and the token is sent back to the client. Whenever a user makes a request from the server, it checks if a token has been included in the request, and checks that token against the list of valid tokens in the 'UserSession' table. If the token exists and is not too old to be used, the user's request is granted, and the user is able to access the resource.

This was effective, but has a few drawbacks. The main drawback is that the token is stored in the database, which means that the database has to be queried every time a user makes a request. This can be slow and can cause performance issues if the database is not optimised. A better implementation of this might have been to use a library that was written to interact on the server fully and store the token in memory, rather than in the database. This would have been faster as querying memory is significantly faster than querying a remote database, as well as allowing the database to not get cluttered with old tokens that have either expired or are no longer in use.

5.3 CRUD Operations

All of the CRUD operations that had to be performed on the data would usually be in response to some event on the front-end of the application on some user interaction. To perform these, it was decided to use Server Actions. These are functions that are defined and designed to be ran on the server, but can be called from anywhere within Next application. The benefit to this is that it meant tat

making direct requests to an endpoint on the server via AJAX or the JS Fetch API is required, and the return types of the functions is known at compile time.

For each of the data models (Course, Module, etc) has their own file in the `‘/actions’` directory, which contains the functions that are used to perform the CRUD operations. The reason for this codebase design was to keep the server actions bundled together, and have each method in a file related to the data that it was retrieving. For the Read operations, the functions could accept a list of strings that could be used to fetch relationship data from the database. This was to allow for the functions to be more diverse.

Each method can also accept a list of the roles that are able to access the data. This was a security feature that was included to ensure that only the correct kind of user can access the data that is being requested. This inclusion ended up not being used however, as role checks we performed on page load itself, so only users that could access a page would trigger the request for the data that would be required to render the page.

5.4 Data Loading

Each of the pages requires data to be retrieved from the database. The method that was used in this project was loading the data on the client side. When a page is mounted to the client (i.e. when the page is loaded), and once the user has been authenticated, the client would make a request to the server to get the data that is required to render the page. This meant that there would be some disconnect between the browser receiving the web page and the page fully rendering. During development, this was not an issue and was overlooked. In the latter stages of the project, it was realised that this was a problem and that the page should be rendered on the server and sent to the client, rather than the other way round.

This is a problem that has not been fixed yet, but is something that would be addressed if there was more time to work on the project. Attempts were made, but many of the implementation methods used for User Authentication and Server Actions were not compatible with the methods used to perform the server-side rendering. This would require a large amount of refactoring to fix, and would be a large task to undertake.

5.5 Component design

For the GoLearn project, the component design was kept simple. Rather than spending an exceptional amount of time developing a complex component library of reusable components, the author decided to use the `‘shadcn/ui’` library. This is an open source library that builds on the `radix-ui` library, and provides a set of components that are designed to be used in React applications. The components are very easy to use and are very flexible, which made them a good choice for the GoLearn project.

5.6 Emails

When a user is created, or they submit a request to reset their password, an email is sent to the user to provide them with a code and a link to set their password. This was implemented using the

react-email and nodemailer libraries. The service that was used to actually send the email was Brevo. This required the setup of a Brevo account, and the configuration of the account in the '.env' file.

5.7 Deployment

The deployment of the GoLearn project was done using Vercel. Vercel is a service that provides hosting for various Node-backed applications, and is very easy to use. The deployment process was very simple once set up. The initial setup required the author to link the Vercel account to the GitHub repository that the project was stored in. The environment variables that were used in the project were then set up in the Vercel dashboard. Once this was done, the author could deploy the project by pushing to a 'master' branch of the repository. This would then trigger a build on Vercel, and the project would be deployed to the Vercel servers.

Once the project was deployed to the Vercel servers, the author then applied a custom domain to the project. This was done by adding the domain to the Vercel dashboard, and then updating the DNS records for the domain to point to the Vercel servers.

5.8 Known Issues

As mentioned before in section 5.4, there is an issue with the way that the pages are rendered. This is related to the the way that the user authentication and data retrieval is currently implemented. At the time of writing this report, the method of loading data onto a page is performed on the client once the page has been sent. When some pages are being loaded, the code to load the data (a 'useEffect' hook) is not ran, which is causing the data to not be loaded, and for the 'loading' state to be stuck on 'true'. To fix this, the pages themselves would need to make use of SSR to validate the auth token and load the data on the server before the page is sent. Then, the data that various components require to run can be passed to specific components as props which can then hold onto that data as state that can then manage the values and make calls back to the server to perform required CRUD operations.

Testing

Unit testing and UAT were both intended to be used for this project. The reason for this was to ensure that the CRUD methods were returning the expected values under different conditions, and that the application was working as expected from a user's perspective.

6.1 Unit Testing

To implement unit testing, the Jest testing framework was used. Jest is a very powerful testing framework that is designed to be used with JS/TS applications. The documentation was referenced to learn how to write tests[16], as well as a number of tutorials that were found online[17][18]. The tests were written to test the Server Actions. The reason that these alone were tested was because they were the only functions that were written that could be tested. The other code that was written was either React components, which have a myriad of problems associated with Unit Testing.

As these server actions were functional wrappers for Prisma queries, the tests needed to also mock out the responses expected from Prisma. This was done by using the 'jest.mock' function to mock out the Prisma functions that were being used in the server actions. This proved difficult to do, largely because there were no video tutorials for doing this, and the documentation was not very clear on how to do it. This page in the documentation was used to implement the mocking of the Prisma functions: <https://www.prisma.io/docs/orm/prisma-client/testing/unit-testing>. This involved setting up a singleton mocking object that could be used to return values that were expected from the Prisma functions.

6.2 UAT

User Acceptance Testing was intended to be performed by the author, as well as by a number of friends and family members. The author would be responsible for testing the application in a number of different scenarios, such as logging in as a student and a teacher, creating and editing content, and enrolling in courses and modules. The friends and family members would also make the same checks, and be watching out for anything that seemed out of place, as well as providing feedback on the design and usability of the application.

GoLearn was not ready for user testing until the week that this report was due. Therefore the

UAT that was performed was very informal, and a questionnaire was not provided for the tester that I managed to get to test the application. That being said, the feedback that was received was very critical and useful in identifying the issues that were present in the application.

The main points that were brought up were as follows:

- “It seems to run fine, a bit odd having the update button be at the top next to ‘user data’ but that might just be my preferences being weird. ”
- “I cant see any courses (did you make a template one?) but from what i can see it all looks good.”
- “I guess you could also move the ‘courses’ from the dropdown to perhaps a bar at the top (as im not sure if its an admin view thing, but the page looks kinda empty)”

These points were very useful in identifying some issues with the UX that seemed like good ideas at the time that they were implemented. The issue mentioned in the second comment was a known issue relating to the way that the data was loaded onto the page. The final comment was something that was considered early on in the project, and was something that was decided against for the purposes of keeping a clean UI. This is something that could be added in the future, but is not a priority for the project.

Some more directed UAT in the form of a questionnaire would be useful in addressing some of the UX concerns brought up by the tester. If the opinions of the tester are shared by other users, then addressing this and making some changes to these aspects of the application would be a good idea.

Legal and Ethical

GDPR[19] refers to the set of regulations surrounding the handling of user data. This set of regulations is designed to protect the data of individuals, and to ensure that companies are handling the data in a responsible and ethical manner. The GoLearn project makes sure that the data that is stored in the database can only be accessed by users with the correct access permissions, and that the data is stored in a secure manner. As well as this, all database communications via prisma are sent via SSL, which means that the data is encrypted when it is sent over the network. This is very important, as it ensures that the data is protected from being intercepted by malicious actors.

User consent The GoLearn project is intended to be used by institutions of learning, such as schools, colleges and universities. This means that the users of the application will already have provided their consent for their data to be stored and maintained and utilised by that institution. The author felt that this would be a suitable level of consent for the application. For this reason, no GDPR consent forms were included in the application, as it was felt that this would be unnecessary.

Project Planning

Atlassian's project management tool, Jira, was used for this project. This was a decision that was decided upon by the University, and was not a decision that was made by the author. Jira provides many features that are useful for project management, such as the ability to create tasks, assign them to team members, and track the progress of the project.

This project makes use of a Kanban Board to track the progress of the project. A Kanban board is a visual representation of the work that needs to be done, and is divided into columns that represent the different stages of the development process. The columns that were used in this project were:

Backlog This column is used to store all of the tasks that need to be done. These tasks are not assigned to anyone, and are not being worked on.

Selected for Development This column is used to store all of the tasks that have been selected for development. These tasks are assigned to a team member, and are expected to be worked on next in the process. This differs from the Backlog column in that the tasks have been selected to be worked on next, rather than being stored for future development.

In Progress This column is used to store all of the tasks that are currently being worked on. These tasks are assigned to a team member, and are being actively worked on.

Done This column is used to store all of the tasks that have been completed. These tasks are assigned to a team member, and are not being actively worked on.

At the beginning of the project, the author created a large number of tasks to be completed, each with various subtasks to bring that task to completion. These tasks were often assigned to a particular story, which is a kind of issue that can be created to represent a use-case scenario that would lead to some feature being implemented. As the project progressed, the author would move tasks from the Backlog column to the Selected for Development column, and then to the In Progress column, and finally to the Done column once the task was completed.

In the first few weeks and months of the project, much work was completed to get the project off the ground and to get the basic functionality working. This involved setting up the NextJS project, creating the data models, setting up the Prisma client, and creating some initial server actions for basic CRUD operations. As the project progressed, the more complex/length features needed to be

implemented. This, coupled with the fact that the author was working on various other work at the same time unrelated to this particular project, meant that the project slowed down significantly. Below (Figure 8.1) is a cumulative flow diagram showing the flow of issues between the various states of completion.




Figure 8.1: Cumulative Flow Diagram

This shows that over time in each term, the beginning would see a peak in the apparent work performed by the author, though this is when the issues were marked as complete as compared to when the work was actually completed.

8.1 GIT

GIT is a VCS method that is a currently one of the most known and used VCS's used. This works by having a set of files and directories being tracked by some program, with changes bundled together in commits. Usually a GIT repository is created and worked on within local environments (on a work laptop or PC for example), and then uploaded to a remote server (called pushing).

For this project, the University provided a GIT remote location (built on top of GitLab). The author also kept track of changes in a GitHub remote on their own personal GitHub account. The University's GitLab setup was to be used and referenced by the University to get statistics on VCS utilisation, whilst the personal remote on GitHub was used as part of a CI/CD with Vercel (as discussed in Section 5.7).

Conclusions

The GoLearn project presented many challenges along the way, ranging from methods of implementing authentication to NextJS specific issues. Many things were learnt by the author along the way. This includes how to write JSX (for React), implementing ORMs and Authentication and GIT for versions control. As these skills had to be learnt in tandem, this led to many refactors of code to get a working solution for the various problems. Overall, despite missing some features that the author wanted from the outset of the project, he is quite happy with the amount of progress that was completed.

9.1 Future Work

There is still a lot of work that needs to go into GoLearn before it can be considered a complete and fully functional application. Some of the work that needs to be done includes:

SSR The first thing that needs to be done is to implement server-side rendering. This is important, as it will allow the pages to be rendered on the server and sent to the client, rather than the other way round. This will help to improve the performance of the application, as well as make it more secure. This is a large task that will require a lot of refactoring, and will take some time to complete.

Assignment Marking and Feedback The next thing that needs to be done is to implement a system for marking assignments and providing feedback to students. This is important, as it will help to ensure that the students are able to learn from their mistakes and improve their work. This is quite a large addition to the application, but would add a lot of value to the project.

Messaging A final feature that needs to be implemented is a messaging system. This would allow for educators and students to communicate with each other about the various modules, units and assignments that have been created and are accessible. This is something that would be completely new to the application, and would only make use of the authentication. Likely, this would require a separate service to run on the server, and have a separate database to store the messages. This is largely because the messages would be so numerous that they would clutter the database that is used to store the other data. A NoSQL database would be a good choice for this, as it would allow for the data to be stored in a more flexible format.

Bibliography

- [1] J. Ferrimanm, "Learndash 3.0." <https://www.learndash.com/best-wordpress-lms-plugin/>, 5 2019.
- [2] Moodle, "Moodle." <https://moodle.org/>, 1 2024.
- [3] Wikipedia and contributors, "Relational database." https://en.wikipedia.org/wiki/Relational_database, 4 2024.
- [4] Wikipedia and contributors, "Nosql." <https://en.wikipedia.org/wiki/NoSQL>, 3 2024.
- [5] M. Wanyoike, "9 best javascript and typescript orms for 2024." <https://www.sitepoint.com/javascript-typescript-orms/>, 3 2023.
- [6] Wikipedia and contributors, "Unit testing." https://en.wikipedia.org/wiki/Unit_testing, 4 2024.
- [7] Wikipedia and contributors, "User acceptance testing." https://en.wikipedia.org/wiki/Usability_testing, 4 2024.
- [8] M. Niels, K. Dempsey, and V. Yan Pilliteri, "An introduction to information security." <https://nvlpubs.nist.gov/nistpubs/SpecialPublications/NIST.SP.800-12r1.pdf>, 4 2024. Section 9.1.4.
- [9] Wikipedia and contributors, "Basic access authentication." https://en.wikipedia.org/wiki/Basic_access_authentication, 4 2024.
- [10] Okta, "What is token-based authentication?." <https://www.okta.com/uk/identity-101/what-is-token-based-authentication/>, 2 2023.
- [11] Okta, "What is token-based authentication?." <https://www.okta.com/uk/identity-101/what-is-token-based-authentication/>, 2 2023. Section 4: JSON Web Token (JWT): A Special Form of Auth Token.
- [12] D. Abramov and R. Nabors, "Introducing react.dev." <https://react.dev/blog/2023/03/16/introducing-react-dev>, 4 2024.
- [13] Next.js, "Server actions and mutations." <https://nextjs.org/docs/app/building-your-application/data-fetching/server-actions-and-mutations>, 4 2024.
- [14] Prisma, "Type safety." <https://www.prisma.io/docs/orm/prisma-client/type-safety>, 4 2024.
- [15] Prisma, "Generating prisma client." <https://www.prisma.io/docs/orm/prisma-client/setup-and-configuration/generating-prisma-client>, 4 2024.

- [16] yepitschunked and Jest, "Getting started." <https://jestjs.io/docs/getting-started>, 1 2024.
- [17] D. Grey, "Next.js with react testing library, jest, typescript." <https://www.youtube.com/watch?v=AS79oJ3Fcf0>, 8 2023.
- [18] M. Maksi, "Practical beginner guide to testing - next.js | react | jest | react testing library (rtl)." <https://www.youtube.com/watch?v=pnLC-9waA44>, 4 2024.
- [19] Gov.uk, "Data protection." <https://www.gov.uk/data-protection>, 4 2024.